

Second order sensitivities in linear or constant time

Roberto Daluise^{*†}

First version: June 23, 2018

This version: November 24, 2018

Abstract

We analyse and compare methods to compute the full set of second order sensitivities of a Monte Carlo price in a time which is at most $O(N \cdot T)$ where N is the number of inputs and T is the time required to compute the price. The new ones include the first algorithm which achieves a complexity $O(T)$ and has acceptable (in fact very low) statistical uncertainties at least in one relevant test case.

Keywords: Greeks; Gamma; Vanna; Volga; Vomma; algorithmic differentiation; adjoints; pathwise differentiation; derivatives pricing; discontinuous pay-offs; digital options.

JEL code: G13.

1 Introduction

In the active management of portfolios of financial products, the sensitivity of the net position to movements of the underlying risk factors is a most precious piece of information. Indeed, one may want to monitor which market shifts could cause a significant jump in the value of the portfolio, and maybe build a hedging portfolio chosen to mitigate the corresponding sensitivities.

In mathematical terms, sensitivities are usually expressed by derivatives of the pricing functional P (also called Greeks). First order Greeks are obviously the most important indicator for a trader to see whether his overall position is approximately market neutral, and are often kept under strict control by dynamically adjusting the hedging positions as time passes. A direct consequence of this behaviour is that the day-to-day performance of the trading desk is at leading order driven by second order Greeks (also called Gammas). By the way, this is the discrete-time counterpart of the fact that in an idealized Black-Scholes world, the continuously rebalanced delta hedging replication strategy generates an Ito component $\frac{1}{2} \frac{\partial^2 P}{\partial X_0^2} dt$ in the dynamics of total wealth.

^{*}Interest Rate and Credit Models, Banca IMI. Address: Largo Mattioli 2, 20100 Milano (Italy). Email: roberto.daluise@bancaimi.com.

[†]Department of Statistics and Quantitative Methods, Milano-Bicocca University. Address: U7 Building, Via Bicocca degli Arcimboldi, 8, 20126 Milano (Italy). Email: r.daluise@campus.unimib.it.

All of the above implies that the knowledge of second order sensitivities is often crucial; which is even more true when the payoff depends on several and maybe correlated risk factors, whose co-movements impact the portfolio value through the mixed terms in the Hessian matrix (sometimes denoted Cross Gammas). This is the case for instance of basket products, indices, multi-asset-class hybrids, and portfolio-level non-linear valuation adjustments such as CVA and DVA (Basel Committee on Banking Supervision, 2015).

Unfortunately, these are exactly the settings in which the pricing exercise is computationally more demanding, as it often involves a multidimensional Monte Carlo simulation. This rules out as infeasible the simplest way to approximate the full $N \times N$ pricing Hessian \mathbf{H} , i.e.

$$H_{ij} \approx \frac{P(\boldsymbol{\psi} + \mathbf{h}_i + \mathbf{h}_j) - P(\boldsymbol{\psi} + \mathbf{h}_i - \mathbf{h}_j) - P(\boldsymbol{\psi} - \mathbf{h}_i + \mathbf{h}_j) + P(\boldsymbol{\psi} - \mathbf{h}_i - \mathbf{h}_j)}{(1 + 3\delta_{ij})h^2},$$

where $\boldsymbol{\psi}$ is the vector of pricing inputs of interest, h is a small displacement and $\mathbf{h}_i := h\mathbf{e}_i$ is h times the i -th element of the canonical basis of \mathbb{R}^N : this method would involve $O(N^2)$ full revaluations.

In fact, the gradient ∇P can be computed in $O(T)$ where T is the time to compute P , thanks to adjoint algorithmic differentiation (AAD, see Giles, 2007; Capriotti, 2011; Capriotti and Giles, 2012), which is more and more widespread for first order sensitivities among practitioners. This technique and its generalizations (e.g. Giles, 2009; Chan and Joshi, 2015; Daluiso and Facchinetti, 2018) coupled with the approximation

$$H_{ij} \approx \frac{\nabla_j P(\boldsymbol{\psi} + h\mathbf{e}_i) - \nabla_j P(\boldsymbol{\psi} - h\mathbf{e}_i)}{2h} \quad \text{or} \quad H_{ij} \approx \frac{\nabla_j P(\boldsymbol{\psi} + h\mathbf{e}_i) - \nabla_j P(\boldsymbol{\psi})}{h},$$

would already give an estimator in $O(N \cdot T)$, but this is in practice quite unsatisfactory, since for small h the result is very noisy, while for moderate h it is too much biased.

The latter fact stimulated research of alternative algorithms which could retain the $O(N \cdot T)$ cost while being unbiased and having less Monte Carlo variance. We point out in particular the significantly different contributions Capriotti (2015), Pagès et al. (2016) and Joshi and Zhu (2016), which seem unaware of one another and present numerical evidences which are difficult to compare.

Our first contribution is therefore a comparative analysis of the three above proposals, to see in which settings they can be applied. In fact, while not delving into the details of the estimators which are already well covered in the original papers, we sometimes go slightly beyond the authors: namely, by underlying a subtlety needed to correctly implement Capriotti (2015), and by developing a multi-fixing generalization of Pagès et al. (2016).

As we will see, in some situations the rich set of alternatives outlined above shrinks considerably: in particular, payoffs directly depending on the non random initial state or on static parameters (e.g. forwards and deterministic interest rates), and factorial models in general where the number of random drivers is smaller than the number of underlyings, are usually out of scope. This motivated the development of the second contribution of this paper, namely a fairly general-purpose

method based on second order pathwise differentiation of the payoff in distributional sense.

Finally, most of the above methods run in $O(N \cdot T)$ time. This is a theoretical upper bound for general computer functions (see e.g. Griewank and Walther, 2008), but can the special structure of a pricing problem with diffusive underlyings change the verdict? To our knowledge, the only partial positive answer is provided by the already cited Joshi and Zhu (2016) in the case in which the simulation can be performed keeping at each time step only a low dimensional state variable.

In this respect, our third and last contribution is a positive answer to the above question in the complementary setting in which the Brownian driver is high dimensional (precisely, has dimensionality at least equal to the number of underlyings). Our solution is based on a relation linking a suitable first order sensitivity to $\partial^2 P / \partial \mathbf{X}_0^2$ via the pricing PDE, and yields the latter in $O(T)$ time by an easily implementable wrapper of any $O(T)$ first order estimator. In fact, also LRMAAD from Capriotti (2015) has constant cost and almost the same perimeter of applicability, but from the numerical analysis it turns out to have too large Monte Carlo uncertainties to be an option in practice; this is sometimes true also for the new estimator, but we will show at least one relevant example in which our new constant-cost algorithm significantly outperforms all its linear-cost competitors.

The rest of the paper is structured as follows. Section 2 states the problem in the diffusive setting in which all our analyses will be performed; Section 3 reviews the methods proposed in the literature; Section 4 introduces the new methods; Section 5 compares the empirical performances and Monte Carlo uncertainties on several test cases; finally, Section 6 summarizes the main findings and concludes.

2 Setting

We consider a description of the market where the risk factors \mathbf{X}_t are modelled as an \mathbb{R}^{N_X} -valued diffusion driven by an N_W -dimensional Wiener process \mathbf{W}_t under a fixed probability measure \mathbb{P} :

$$d\mathbf{X}_t = \boldsymbol{\mu}(\boldsymbol{\theta}, \mathbf{X}_t, t)dt + \boldsymbol{\Sigma}(\boldsymbol{\theta}, \mathbf{X}_t, t)d\mathbf{W}_t, \quad t \geq 0, \quad (2.1)$$

where $\boldsymbol{\theta} \in \mathbb{R}^{N_\theta}$ is a vector of parameters. A price P of interest is expressed as an expected value

$$P = \mathbb{E}[g(\mathbf{X}_{T_1}, \dots, \mathbf{X}_{T_M}, \boldsymbol{\theta})] =: \mathbb{E}[g(\mathbf{X}_T, \boldsymbol{\theta})], \quad (2.2)$$

and is estimated by Monte Carlo simulation; to be concrete, we suppose that the simulation is performed by Euler stepping on a time grid $\{S_0, \dots, S_L\} = \mathbf{S} \supseteq \mathbf{T}$:

$$\begin{aligned} \mathbf{X}_{S_{i+1}} &= \mathbf{X}_{S_i} + \boldsymbol{\mu}(\boldsymbol{\theta}, \mathbf{X}_{S_i}, S_i) \Delta_i S + \boldsymbol{\Sigma}(\boldsymbol{\theta}, \mathbf{X}_{S_i}, S_i) \sqrt{\Delta_i S} \cdot \mathbf{Z}^{(i)}, \\ \Delta_i S &:= S_{i+1} - S_i, \quad \mathbf{Z}^{(i)} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_d), \end{aligned} \quad (2.3)$$

although one could imagine more or less straightforward generalizations of the methods described below for more general schemes.

When the payoff does not depend on $\boldsymbol{\theta}$ directly, we will write

$$g(\mathbf{X}_T, \boldsymbol{\theta}) =: p(\mathbf{X}_T); \quad (2.4)$$

Name	1st order	2nd order	Section	Reference
FDIFF2	finite difference	finite difference	1	N.A.
FDPW	pathwise AD	finite difference	3.1	N.A.
LRMAAD	pathwise AD	likelihood ratio	3.2	Capriotti (2015)
AADLRM	likelihood ratio	pathwise AD	3.2	Capriotti (2015)
VAD	vibrato	pathwise AD	3.3	Pagès et al. (2016)
VFD	vibrato	finite difference	3.3	N.A.
HOPP	change var + AD	pathwise AD	3.4	Joshi and Zhu (2016)

Table 1: Legacy methods to compute second order sensitivities: estimators used for first and second order differentiation, references. Here and elsewhere in the paper, AD stands for algorithmic differentiation.

we will call “autonomous” this simpler case. Note that the price P still depends on θ through the dynamics (2.1).

We aim at computing the full Gamma matrix

$$\mathbf{\Gamma} := \frac{\partial^2 P}{\partial \mathbf{X}_0^2},$$

and possibly the full Hessian

$$\mathbf{H} := \frac{\partial^2 P}{\partial \psi^2}, \quad \psi := (\mathbf{X}_0, \theta),$$

in acceptable time when N_X and N_θ are large. Note that a naive approach based on finite differences would involve

$$O\left((N_X + N_\theta)^2\right) =: O(N^2)$$

re-evaluations of the price P with displaced values of θ and/or \mathbf{X}_0 .

Throughout the paper, the following notational conventions are adopted: scalars are in *italic* (e.g. a), vectors in **boldface italic** (e.g. \mathbf{a} has components a_i), matrices in straight **boldface** (e.g. \mathbf{A} has entries A_{ij}); vectors are interpreted as columns unless transposed; for a vector \mathbf{a} , the gradient $\bar{\mathbf{a}} = \frac{\partial P}{\partial \mathbf{a}}$ is a row vector; for a matrix \mathbf{A} , the gradient $\bar{\mathbf{A}} = \frac{\partial P}{\partial \mathbf{A}}$ is a matrix with components $\bar{A}_{ij} = \frac{\partial P}{\partial A_{ji}}$ (note the inversion of indices). Finally, for stochastic processes, we will write the time argument in the subscript (e.g. \mathbf{X}_t) if no confusion arises, or as a function argument (e.g. $X_i(t)$) otherwise.

3 Review of existing algorithms

In this section, we review from an algorithmic point of view the main methods based on the existing literature whose theoretical complexity is $O(N \cdot T)$, with a focus on the perimeter of applicability within the setting of Section 2. For quick reference, we gather in Table 1 the essential traits of all these algorithms.

3.1 Finite differences of pathwise adjoint

Before describing more complex methods, let us mention that the computation of finite differences of the pathwise adjoint estimator is already a linear cost method. Indeed, first order pathwise adjoints cost at most $4T$; repeating their computation N times for small directional displacements of ψ , one trivially gets an estimator for the Hessian, which will be denoted by the acronym FDPW.

However, we expect poor variance properties of this naive method, as finite differences have high Monte Carlo uncertainties when the integrand is discontinuous, as is the case for the gradient of typical payoffs. Therefore, we explore in the following sections analytical alternatives leveraging on the structure of the problem.

3.2 Likelihood ratio and adjoint differentiation

The idea of Capriotti (2015) is to combine pathwise differentiation with the likelihood ratio method. Since the latter can handle only the autonomous case, the same limitation applies to the resulting second order methods; moreover, as is clear from the matrix inversion in equation (3.2) here below, one needs that $N_X \leq N_W$ with $\Sigma(\theta, \mathbf{X}_T, T)$ almost surely full rank. As we believe that a crucial detail in the correct interpretation of the final formula may be easily overlooked, we repeat the derivation briefly.

When adjoint algorithmic differentiation is applied first, one wants to compute

$$\frac{\partial}{\partial \psi} \mathbb{E} \left[\frac{\partial p}{\partial \mathbf{X}_T}(\mathbf{X}_T) \frac{\partial \mathbf{X}_T}{\partial \psi}(\mathbf{Z}, \psi) \right]$$

by likelihood ratio. In order to do so, one *must* rewrite the argument of \mathbb{E} so that the only random variable appearing is \mathbf{X}_S , because the method relies on writing the expectation as an integral against the density of \mathbf{X}_S . Fortunately, $\mathbf{Z}^{(i)}$ is readily expressed as a function of \mathbf{X}_S and θ as

$$z^{(i)}(\mathbf{X}_S, \theta) = \Sigma(\theta, \mathbf{X}_{S_i}, S_i)^{-1} (\Delta_i S)^{-\frac{1}{2}} (\mathbf{X}_{S_{i+1}} - \mathbf{X}_{S_i} - \boldsymbol{\mu}(\theta, \mathbf{X}_{S_i}, S_i) \Delta_i S).$$

Now a straightforward generalization of the likelihood ratio method to the case in which the integrand has explicit dependence on the differentiation parameter leads to the following estimator of the Hessian \mathbf{H} :

$$\left(\frac{\partial p}{\partial \mathbf{X}_T}(\mathbf{X}_T) \frac{\partial \mathbf{X}_T}{\partial \psi}(\mathbf{Z}, \psi) \right)^\top \boldsymbol{\Omega}_\psi + \frac{\partial}{\partial \psi} \left[\frac{\partial p}{\partial \mathbf{X}_T}(\mathbf{X}_T) \frac{\partial \mathbf{X}_T}{\partial \psi}(z(\mathbf{X}_S, \theta), \psi) \right], \quad (3.1)$$

where the $\boldsymbol{\Omega}_\psi$ is the customary likelihood ratio weight capturing the component of the derivative due to dependence of the distribution of \mathbf{X}_S on ψ by differentiation of the log-density, while the future path $\mathbf{X}_{S \setminus \{0\}}$ in the square bracket is a constant for the purpose of differentiation. Note that in general the second addend cannot be replaced by the simpler

$$\frac{\partial p}{\partial \mathbf{X}_T}(\mathbf{X}_T) \frac{\partial}{\partial \psi} \left[\frac{\partial \mathbf{X}_T}{\partial \psi}(\mathbf{Z}, \psi) \right]$$

which seems implicit in the implementation described by Capriotti (2015).

Therefore, the algorithm is as follows:

Algorithm 3.1 (LRMAAD).

1. Compute simultaneously all $\bar{\mathbf{X}}_{T_i} = \frac{\partial p}{\partial \mathbf{X}_{T_i}}$ with AAD.
2. Compute the derivatives of the Gaussian transitions' log-densities as

$$\begin{aligned}\bar{\boldsymbol{\mu}}^{(i)} &= \left(\mathbf{X}_{S_{i+1}} - \boldsymbol{\mu}^{(i)} \right)^\top V_i^{-1}, \\ \bar{\boldsymbol{\Sigma}}^{(i)} &= \left(\boldsymbol{\Sigma}^{(i)} \right)^\top \left[\left(\bar{\boldsymbol{\mu}}^{(i)} \right)^\top \left(\bar{\boldsymbol{\mu}}^{(i)} \right) - V_i^{-1} \right],\end{aligned}\tag{3.2}$$

where the mean, standard deviation and covariance matrices are defined as

$$\begin{aligned}\boldsymbol{\mu}^{(i)} &= \mathbf{X}_{S_i} + \boldsymbol{\mu}(\boldsymbol{\theta}, \mathbf{X}_{S_i}, S_i) \Delta_i S, \\ \boldsymbol{\Sigma}^{(i)} &= \boldsymbol{\Sigma}(\boldsymbol{\theta}, \mathbf{X}_{S_i}, S_i) \sqrt{\Delta_i S}, \quad V_i = \left(\boldsymbol{\Sigma}^{(i)} \right) \left(\boldsymbol{\Sigma}^{(i)} \right)^\top.\end{aligned}$$

3. Propagate them backward with AAD to $\boldsymbol{\psi}$ obtaining likelihood ratio weights Ω_j for $j = 1, \dots, N$.
4. Propagate $\bar{\mathbf{X}}_{T_i}$ backward through simulation obtaining $\bar{\boldsymbol{\psi}} = \sum_{i=1}^M \bar{\mathbf{X}}_{T_i} \cdot \frac{\partial \mathbf{X}_{T_i}}{\partial \boldsymbol{\psi}}$.
5. Apply algorithmic differentiation or finite differences to the previous step, where $\bar{\mathbf{X}}_{T_i}$ should be considered as constants while $\mathbf{Z}^{(i)}$ should be *recomputed* as

$$\left(\boldsymbol{\Sigma}^{(i)} \right)^{-1} \left(\mathbf{X}_{S_{i+1}} - \boldsymbol{\mu}^{(i)} \right).$$

Call \mathbf{J} the resulting Jacobian with respect to $\boldsymbol{\psi}$.

6. The estimator for H_{ij} is $J_{ij} + \bar{\psi}_i \Omega_j$.

Only step 5 has a cost linearly growing with N . This is unavoidable in general, since computing a Jacobian \mathbf{J} multiplies the run time of a function by $O(N)$, regardless of whether one uses AAD as suggested in the original article, or a simpler forward differentiation (Griewank and Walther, 2008), or even finite differences.

However, if one wants to determine only $\boldsymbol{\Gamma}$, then this term can be computed very efficiently. Indeed, the step under analysis computes the Jacobian of the second addend of (3.1), which for $\boldsymbol{\psi} = \mathbf{X}_0$ reduces to

$$\frac{\partial}{\partial \mathbf{X}_0} \left[\frac{\partial p}{\partial \mathbf{X}_T}(\mathbf{X}_T) \frac{\partial \mathbf{X}_T}{\partial \mathbf{X}_0} (z(\mathbf{X}_S), \mathbf{X}_0) \right];$$

now, one can easily note that if the quantity in square brackets is computed by AAD, then in the implementation it takes the form

$$\bar{\mathbf{X}}_{S_1} \cdot \frac{\partial \mathbf{X}_{S_1}}{\partial \mathbf{X}_0} (z(\mathbf{X}_{S_1}, \mathbf{X}_0), \mathbf{X}_0)$$

for a suitable random variable $\bar{\mathbf{X}}_{S_1}$ *not depending on* \mathbf{X}_0 . Once this is remarked, differentiation of this term is so fast to become practically negligible, and the complexity of the algorithm is dominated by the first term in (3.1), computed in $O(T)$

time. Note that even there, the likelihood ratio weight $\Omega_{\mathbf{X}_0}$ has non-trivial contribution from (3.2) only for the first step $i = 0$: this makes the overall runtime essentially comparable to that of computing the first order pathwise estimator

$$\frac{\partial p}{\partial \mathbf{X}_T}(\mathbf{X}_T) \frac{\partial \mathbf{X}_T}{\partial \mathbf{X}_0}(\mathbf{Z}, \mathbf{X}_0).$$

Instead, if one differentiates pathwise the first order likelihood ratio estimator, the following is obtained:

Algorithm 3.2 (AADLRM). Implement Algorithm 3.1 replacing the steps:

5. Apply algorithmic differentiation (forward or backward) or finite differences to the computation of Ω (i.e. of the simulation scheme and of steps 2 and 3), obtaining its Jacobian \mathbf{J} with respect to ψ .
6. The estimator for H_{ij} is $J_{ij}p(\mathbf{X}_T) + \bar{\psi}_j \Omega_i$.

Again the Jacobian step 5 is a potential $O(N)$ factor on run times. The original paper notes that this second method is empirically slower on a simple example, maybe because of the above Γ -specific speedup; however, we feel that a comparison of Monte Carlo variances is needed before discarding AADLRM as less performing.

3.3 Vibrato and adjoint differentiation

Pagès et al. (2016) apply automatic differentiation to a different first order estimator called Vibrato Monte Carlo, which has much better sample variance properties especially after the introduction of a trick based on the antithetic transform. (They also consider the possibility of applying Vibrato twice, but they find no relevant difference.) As in the previous section, the analysis works only for the autonomous case and under the hypothesis $N_X \leq N_W$ with $\Sigma(\theta, \mathbf{X}_T, T)$ almost surely full rank.

The algorithm to compute the estimator, here generalized to $M > 1$ fixing times in the spirit of what Giles (2009) mentions for first order Vibrato, is as follows:

Algorithm 3.3 (VAD).¹

1. For each fixing time $T_i = S_{t(i)+1}$, compute the derivatives of the (approximate) Gaussian log-density of \mathbf{X}_{T_i} given its neighbours in the simulation grid \mathbf{S} : i.e. apply (3.2) where

$$\begin{aligned} \mu^{(N)} &= \mathbf{X}_{S_{t(N)}} + \mu(\theta, \mathbf{X}_{S_{t(N)}}, S_{t(N)}) \Delta_{t(N)} S, \\ \Sigma^{(N)} &= \Sigma(\theta, \mathbf{X}_{S_{t(N)}}, S_{t(N)}) \sqrt{\Delta_{t(N)}} S, \\ \mu^{(i)} &= (1 - \lambda_{t(i)}) \mathbf{X}_{S_{t(i)}} + \lambda_{t(i)} \mathbf{X}_{S_{t(i)+2}} \quad \text{for } i < N, \\ \Sigma^{(i)} &= \sqrt{(\Delta_{t(i)} S) (1 - \lambda_{t(i)})} \Sigma(\theta, \mathbf{X}_{S_{t(i)}}, S_{t(i)}) \quad \text{for } i < N, \\ \lambda_t &:= (\Delta_t S + \Delta_{t+1} S)^{-1} \Delta_t S \quad \text{and} \quad V_i := \left(\Sigma^{(i)} \right) \left(\Sigma^{(i)} \right)^\top \quad \text{for } i \leq N. \end{aligned} \tag{3.3}$$

¹The algorithms in the previous section are called LRMAAD when LRM is applied *after* AAD, and AADLRM otherwise; while here, somewhat confusingly, VAD is the method which applies V(ibrato) *before* AD. We keep the acronyms chosen by the original authors nonetheless.

2. For each fixing time $T_i = S_{t(i)+1}$, draw \mathbf{X}_{T_i} from the conditional law:

$$\mathbf{X}_{T_i} = \boldsymbol{\mu}^{(i)} + \boldsymbol{\Sigma}^{(i)} \mathbf{W}^{(i)}, \quad \mathbf{W}^{(i)} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_{N_W}) \quad (3.4)$$

3. Compute the antithetic variates

$$p_i^- := p \left(X_{T_1}, \dots, \boldsymbol{\mu}^{(i)} - \boldsymbol{\Sigma}^{(i)} \mathbf{W}^{(i)}, \dots, X_{T_M} \right),$$

$$p_i^0 := p \left(X_{T_1}, \dots, \boldsymbol{\mu}^{(i)}, \dots, X_{T_M} \right).$$

4. Multiply $\bar{\boldsymbol{\mu}}^{(i)}$ by $\frac{p - p_i^-}{2}$ and $\bar{\boldsymbol{\Sigma}}^{(i)}$ by $\frac{p + p_i^-}{2} - p_i^0$.
5. Repeat steps 2 to 4 for N_{in} independent draws of the random numbers \mathbf{W} and average the resulting $\bar{\boldsymbol{\mu}}^{(i)}$, $\bar{\boldsymbol{\Sigma}}^{(i)}$.
6. Propagate backwards $\bar{\boldsymbol{\mu}}^{(i)}$, $\bar{\boldsymbol{\Sigma}}^{(i)}$ through the adjoint of (3.3) and of the Euler scheme, obtaining sensitivities $\bar{\boldsymbol{\psi}}$.
7. Apply automatic differentiation (forward or backward) to the whole algorithm.

Note that the last step multiplies by $O(N)$ the cost of the whole algorithm. As for the methods of the previous section, if the payoff is discontinuous then the method cannot be applied as is; however, in this case we can conceive to perform the last step with a finite difference with non-small displacement: we will denote this variant as VFD.

Finally, the choice of N_{in} would need investigation, since Pagès et al. (2016) suggest a small N_{in} (say 1 or 2), while Giles (2009) finds that at least for first order sensitivities a good value of N_{in} from the variance point of view should be of the order of the square root of the outer Monte Carlo scenarios. However, when coupled with the antithetic technique as generalized above for $M > 1$ fixings, the latter advice would result in a complexity per outer scenario which is much higher than in the original pricing, and this may be a sufficient concern in applications to give up the improvement in asymptotic uncertainties.

3.4 Optimal partial proxy for Hessians

In real payoffs, direct pathwise second order differentiation is not possible only because of a few discontinuity points of the gradient or of the payoff. Exploiting this fact,

Joshi and Zhu (2016) manage to change the integrand without changing the integral, so that one can differentiate algorithmically twice the modified payoff. More precisely, they rewrite the expected value as a function of a multivariate uniform \mathbf{U} ,

$$P = \mathbb{E} [g(\mathbf{X}_T, \boldsymbol{\theta})] =: \mathbb{E} [g(\mathbf{x}_T(\mathbf{Z}, \boldsymbol{\psi}), \boldsymbol{\theta})] =: \mathbb{E} [g(\mathbf{x}_T(\mathcal{N}^{-1}(\mathbf{U}), \boldsymbol{\psi}), \boldsymbol{\theta})]$$

and then perform a change of variables $\mathbf{V} = \mathbf{v}(\mathbf{U}, \boldsymbol{\psi})$, designed in such a way that while moving $\boldsymbol{\psi}$, the path generated by $\mathbf{v}(\mathbf{U}, \boldsymbol{\psi})$ does not cross the discontinuities. After the change of variables,

$$P = \mathbb{E} [g(\mathbf{x}_T(\mathcal{N}^{-1}(\mathbf{v}(\mathbf{U}, \boldsymbol{\psi})), \boldsymbol{\psi}), \boldsymbol{\theta}) \cdot \omega(\mathbf{U}, \boldsymbol{\psi})]$$

for a suitable weight ω ; the mapping \mathbf{v} should ensure that if the payoff is smooth except when

$$U_{N_W}^{(t(i))} = a_j^{(i)} \left(\boldsymbol{\psi}, \mathbf{U}^{(1)}, \dots, \mathbf{U}^{(t(i)-1)}, U_1^{(t(i))}, \dots, U_{N_W-1}^{(t(i))} \right) \quad \text{for some } j,$$

where $\{T_{t(i)+1}\}_{i=1,\dots,h}$ are the maturities of the payoff's discontinuities, then²

$$U_{N_W}^{(t(i))} \in \left[a_j^{(i)}(\boldsymbol{\psi}, \mathbf{U}), a_{j+1}^{(i)}(\boldsymbol{\psi}, \mathbf{U}) \right) \iff \forall \boldsymbol{\psi}, V_{N_W}^{(t(i))} \in \left[a_j^{(i)}(\boldsymbol{\psi}, \mathbf{V}), a_{j+1}^{(i)}(\boldsymbol{\psi}, \mathbf{V}) \right).$$

The simplest way to build such \mathbf{v} is to make $\mathbf{v}^{(t)}$ depend only on $(\mathbf{U}^{(s)})_{s \leq t}$ and keep it equal to the identity mapping except on coordinates $(v_{N_W}^{(t(i))})_{i \geq 1}$; this way the change-of-variables weight is

$$\omega(\mathbf{U}, \boldsymbol{\psi}) = \prod_{i \geq 1} \frac{\partial v_{N_W}^{(t(i))}}{\partial U_{N_W}^{(t(i))}}.$$

The remaining degrees of freedom are fixed requiring that $\mathbf{v}(\mathbf{U}, \boldsymbol{\psi}_0) = \mathbf{U}$ for a reference value $\boldsymbol{\psi}_0$ which is the point where ∇P is to be estimated, and then minimizing the conditional variance of $\partial\omega/\partial\boldsymbol{\psi}$: it turns out that the best \mathbf{v} in this sense is obtained when $V_{N_W}^{(t(i))}$ linearly interpolates between $a_j^{(i)}$ and $a_{j+1}^{(i)}$.³

Joshi and Zhu (2016) proceed suggesting an implementation based on a smart backpropagation of second order derivatives (Joshi and Yang, 2011), which should improve the computational burden when there are significantly fewer state variables in the stochastic process than parameters in the model. We do not pursue this level of optimization in our implementation, which aims at being generic; the analysis of the applicability and efficacy of such model-dependent and payoff-dependent trick on the different competitors described in this paper could be the subject of future research.⁴ For the moment being, we sketch a less *ad hoc* algorithm to compute the estimator:

Algorithm 3.4 (HOPP).

1. During simulation, compute the critical values $a_j^{(i)}$, the transformed drivers \mathbf{V} and the weight ω . (In fact we know *a priori* that $\mathbf{V} = \mathbf{U}$ and $\omega = 1$ at $\boldsymbol{\psi} = \boldsymbol{\psi}_0$, but the instructions computing \mathbf{V} and ω should be taken into account when in the sequel this step will be differentiated.)
2. Multiply the payoff g by the weight ω (with the same remark about $\omega = 1$).
3. Differentiate algorithmically twice the previous two steps, inclusive of simulation, with respect to $\boldsymbol{\psi}$ for fixed $\boldsymbol{\psi}_0$.

Step 3 is an $O(N)$ multiplier on the execution time of the first two steps, which in turn is typically little more than that of a plain Monte Carlo run, the overhead of computing ω being slightly significant only when the discontinuities are very frequent as in barrier options.

²At least to second order, as the original article points out.

³In fact, between their second order Taylor expansion in $\boldsymbol{\psi}$ around $\boldsymbol{\psi}_0$.

⁴Note that the optimization would not apply anyway to the numerical tests of this paper, whose validity is therefore unaffected by the simplification.

Name	1st order	2nd order	Section
FDDAAD	distributional AD	finite differences	4.2
DAAD2	(distributional) AD	distributional AD	4.2
FT	any	none	4.3

Table 2: New methods to compute second order sensitivities: estimators used for first and second order differentiation, references. See also Table 1 for legacy methods.

4 New methods

In this section we move to new proposals which try to overcome some limitations of the above methods. Table 2 completes with the contents of this section the summary which was provided in Table 1 for the algorithms of Section 3.

4.1 Second order differentiation by conditioning

We isolate in this small section a simple mathematical property, which will be used to derive the new second order method of Section 4.2, but which is meaningful *per se*. The general question is whether the ability to compute efficiently an estimator of the Hessian of a *conditional* expectation can enable the fast computation of the Hessian of the unconditional expectation.

More specifically, we take a set of times $\tilde{T} \subseteq S$, and let

$$\varepsilon(\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta}) = \mathbb{E}[g(\mathbf{X}_T, \boldsymbol{\theta}) | \mathbf{X}_{\tilde{T}}].$$

Then one can compute the desired Hessian as $\mathbb{E}\left[\frac{\partial^2}{\partial \boldsymbol{\psi}^2} \varepsilon(\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta})\right]$, i.e.

$$\mathbb{E}\left[\left(\frac{\partial(\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta})}{\partial \boldsymbol{\psi}}\right)^\top \frac{\partial^2 \varepsilon}{\partial (\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta})^2} \left(\frac{\partial(\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta})}{\partial \boldsymbol{\psi}}\right) + \frac{\partial \varepsilon}{\partial \mathbf{X}_{\tilde{T}}} \left(\frac{\partial^2 \mathbf{X}_{\tilde{T}}}{\partial \boldsymbol{\psi}^2}\right)\right], \quad (4.1)$$

where by the tower rule, one can substitute the terms

$$\frac{\partial \varepsilon}{\partial \mathbf{X}_{\tilde{T}}}, \frac{\partial^2 \varepsilon}{\partial (\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta})^2}$$

with $\mathbf{X}_{\tilde{T}}$ -conditional sample estimators. We note that the second addend of (4.1) is readily computed in linear time by differentiation of the function

$$\boldsymbol{\psi} \mapsto \bar{\mathbf{X}}_{\tilde{T}} \cdot \frac{\partial \mathbf{X}_{\tilde{T}}}{\partial \boldsymbol{\psi}}, \quad \text{with} \quad \bar{\mathbf{X}}_{\tilde{T}} \text{ any estimator of } \frac{\partial \varepsilon}{\partial \mathbf{X}_{\tilde{T}}},$$

which function is easily computed by pathwise AAD; the Jacobians in the first addend of (4.1) are similarly not a problem, and so we just need an estimator of the Hessian of ε , as desired.

4.2 Distributional algorithmic differentiation

The methods of Section 3 are devised to avoid at least in part the need of second order differentiation of the payoff, because its first derivative is typically discontinuous: a vanilla call already has

$$\frac{d}{dX}(X - K)^+ = I_{X > K}.$$

However, Daluiso and Facchinetti (2018) present a generalization of the first order pathwise method which works for discontinuous payoffs. Then on the one hand, by finite differences one may immediately form a simple biased second order estimator for discontinuous payoffs (FDDAAD); on the other hand, a new unbiased method for both continuous and discontinuous payoffs working also in the non-autonomous case can be introduced as follows.

As in most applications, suppose that there exist smooth scalar functions

$$f_i(\mathbf{X}_T, \boldsymbol{\theta}), g_a(\mathbf{X}_T, \boldsymbol{\theta}) \quad \text{for} \quad i \in 1, \dots, h, \mathbf{a} \in \{-1, +1\}^h,$$

such that g is naturally represented as

$$g(\mathbf{X}_T, \boldsymbol{\theta}) = g_a(\mathbf{X}_T, \boldsymbol{\theta}) \quad \text{on } \mathcal{X}_a := \{(\mathbf{X}_T, \boldsymbol{\theta}) : a_i \cdot f_i(\mathbf{X}_T, \boldsymbol{\theta}) \geq 0 \forall i = 1, \dots, h\}.$$

We call $M_i = S_{t(i)+1}$ the last time in \mathbf{T} such that f_i depends on \mathbf{X}_{M_i} ; we suppose that we can find $k(i)$ such that the $k(i)$ -th component of

$$\frac{\partial f_i}{\partial \mathbf{Z}^{(t(i))}} = \frac{\partial f_i}{\partial \mathbf{X}_{M_i}} \boldsymbol{\Sigma}^{(t(i))}$$

is almost-surely different from zero.

Now we exploit the conditioning idea of (4.1) using as $\tilde{\mathbf{T}}$ the (often quite small) set obtained substituting in \mathbf{T} each M_i with its first neighbours in the simulation grid \mathbf{S} . Recall that \mathbf{X}_{M_i} can be simulated conditionally on its neighbours by Brownian bridge interpolation (3.4). With this choice of $\tilde{\mathbf{T}}$, we can estimate the Hessian of ε by applying first pathwise differentiation to its definition (with distributional corrections if the payoff itself is discontinuous), and then distributional algorithmic differentiation to the result.

A brief derivation of the full estimator for possibly discontinuous payoffs is available in Appendix A; here, for the sake of clarity, we detail the algorithm under the hypothesis that only the gradient is discontinuous.

Call \tilde{f}_i, \tilde{g} the functions computing f_i, g from $\mathbf{X}_{\tilde{\mathbf{T}}}$, including conditional simulation of \mathbf{X}_{M_i} . Then the algorithm on each Monte Carlo scenario is as follows:

Algorithm 4.1. [DAAD2]

1. Apply pathwise adjoint differentiation to the composition of g with the conditional simulation of the \mathbf{X}_{M_i} , getting estimators $\bar{\boldsymbol{\theta}}, \bar{\mathbf{X}}_{\tilde{\mathbf{T}}}$.
2. Differentiate again the previous step ignoring discontinuities, getting a matrix \mathbf{D}_0 representing the smooth part of the $\mathbf{X}_{\tilde{\mathbf{T}}}$ -conditional sample estimator for the Hessian of ε .

3. For each $i = 1, \dots, h$ compute the values ω which substituted into $W_{k(i)}^{(i)}$ make a $\tilde{\mathbf{W}}^{(i)}$ such that f_i equals zero, and the the corresponding perturbed version of \mathbf{X}_T , which we denote by $\tilde{\mathbf{X}}_T$. Then for each such ω :

- (a) There should be exactly two $\mathbf{a} \in \{-1, +1\}^h$ such that $\tilde{\mathbf{X}}_T \in \mathcal{X}_{\mathbf{a}}$, corresponding to $a_i = \pm 1$: call them \mathbf{a}_{\pm} .
- (b) Use AAD to compute in the perturbed scenario the gradients

$$\frac{\partial \tilde{f}_i}{\partial \mathbf{W}^{(i)}}, \frac{\partial \tilde{f}_i}{\partial \mathbf{X}_{\tilde{T}}}, \frac{\partial \tilde{f}_i}{\partial \boldsymbol{\theta}}, \frac{\partial \tilde{g}_{\mathbf{a}_{\pm}}}{\partial \mathbf{X}_{\tilde{T}}}, \frac{\partial \tilde{g}_{\mathbf{a}_{\pm}}}{\partial \boldsymbol{\theta}}.$$

- (c) Increment the discontinuity's contribution to $\frac{\partial^2 \varepsilon}{\partial (\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta})^2}$ as follows:

$$\mathbf{D}_i += \frac{e^{-\frac{\omega^2}{2}}}{\sqrt{2\pi}} \left| \frac{\partial \tilde{f}_i}{\partial W_{k(i)}^{(i)}} \right|^{-1} \left(\frac{\partial \tilde{f}_i}{\partial (\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta})} \right)^{\top} \left[\frac{\partial g_{\mathbf{a}_{+}}}{\partial (\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta})} - \frac{\partial g_{\mathbf{a}_{-}}}{\partial (\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta})} \right]. \quad (4.2)$$

4. Compute the Jacobian $\frac{\partial \mathbf{X}_{\tilde{T}}}{\partial \boldsymbol{\psi}}$, which is the only unknown part of $\frac{\partial (\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta})}{\partial \boldsymbol{\psi}}$.
5. Compute by adjoint differentiation $\bar{\mathbf{X}}_{\tilde{T}} \cdot \frac{\partial \mathbf{X}_{\tilde{T}}}{\partial \boldsymbol{\psi}}$, then differentiate it again with respect to $\boldsymbol{\psi}$ getting a matrix $\bar{\bar{\boldsymbol{\psi}}}$.
6. The final estimator is $\bar{\bar{\boldsymbol{\psi}}} + \left(\frac{\partial (\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta})}{\partial \boldsymbol{\psi}} \right)^{\top} \left(\sum_{i=0}^h \mathbf{D}_i \right) \left(\frac{\partial (\mathbf{X}_{\tilde{T}}, \boldsymbol{\theta})}{\partial \boldsymbol{\psi}} \right)$.

As for complexity, the multiplier on the cost of simulation comes from step 5, with a factor of order $O(N)$, plus step 4, with a factor either of order $O(h \cdot N_X)$ if computed by adjoint differentiation, or of order $O(N)$ if computed by forward differentiation. The cost of steps 1 through 3 has to do with the number of discontinuities h , but is also proportional to the computational time of \tilde{f}_i and \tilde{g} , which is usually much less relevant than the time required to perform the Euler simulation.

4.3 Functional Gamma

The idea of this section is that if $N_X \leq N_W$, the Feynman-Kac theorem can supply enough equations to determine the Gamma matrix $\boldsymbol{\Gamma} = \partial^2 P / \partial \mathbf{X}_0^2$ given the time decay of suitable first order sensitivities, whose computation is fast ($O(T)$) thanks to adjoint differentiation. This fact is essentially a consequence of the multi-dimensional version of a known relationship between Gamma and “functional Vega”, which is exploited for instance in Reghai et al. (2015) in a one-dimensional setting for purposes not related to the computation of a high number of sensitivities.

Indeed, it is well known that the price P , as a function of \mathbf{X}_0 and of the evaluation time, satisfies the PDE

$$\frac{\partial P}{\partial t} + \frac{\partial P}{\partial x} \cdot \boldsymbol{\mu} + \frac{1}{2} \text{tr} \left[\frac{\partial^2 P}{\partial x^2} \boldsymbol{\Sigma} \boldsymbol{\Sigma}^{\top} \right] - Pr + \sum C_i \delta(t - t_i) = 0,$$

where r is the risk-free rate and C_i is the cashflow paid at time t_i . Now one can note that if P' is the price for a different diffusion coefficient function Σ' , then $P' - P$ satisfies the same equation except that the discrete last addend is substituted by a continuous source term

$$\frac{1}{2} \text{tr} \left[\frac{\partial^2 P}{\partial x^2} (\Sigma' \Sigma'^\top - \Sigma \Sigma^\top) \right],$$

so that by Feynman-Kac we can conclude that

$$P'(x, t) = P(x, t) + \frac{1}{2} \mathbb{E} \left\{ \int_t^T e^{-\int_t^s r(\mathbf{X}_u, u) du} \text{tr} \left[\frac{\partial^2 P}{\partial x^2} (\Sigma' \Sigma'^\top - \Sigma \Sigma^\top) \right] (\mathbf{X}_s, s) ds \right\}.$$

Now we fix a perturbation matrix \mathbf{H} and time $\tau > 0$ and choose

$$\Sigma'(t, x) = \Sigma(t, x) + \mathbf{H} I_{t < \tau};$$

differentiation with respect to \mathbf{H} gives

$$\frac{\partial P'(x, 0)}{\partial \mathbf{H}} = \int_0^\tau \mathbb{E} \left\{ e^{-\int_0^s r(\mathbf{X}_u, u) du} \left[\Sigma^\top \frac{\partial^2 P}{\partial x^2} \right] (\mathbf{X}_s, s) \right\} ds;$$

finally, differentiation in $\tau = 0$ gives

$$\frac{\partial^2 P'(x, 0)}{\partial \tau \partial \mathbf{H}} = \Sigma^\top (\mathbf{X}_0, 0) \mathbf{\Gamma}, \quad (4.3)$$

which for full rank $\Sigma(\mathbf{X}_0, 0)$ allows for the determination of $\mathbf{\Gamma}$ given $\partial^2 P'(x, 0)/(\partial \tau \partial \mathbf{H})$; we expect the computation of the latter to be faster because the second order differentiation is with respect to a one dimensional variable τ .

To compute $\partial^2 P'(x, 0)/(\partial \tau \partial \mathbf{H})$, we note that for small τ , the payoff depends on \mathbf{H} only through the first simulated point \mathbf{X}_{S_1} . This makes the following decomposition possible:

$$\frac{\partial P'}{\partial \mathbf{H}} = \mathbb{E} \left[\frac{\partial P(\mathbf{X}_{S_1}, S_1)}{\partial \mathbf{X}_{S_1}} \cdot \frac{\partial \mathbf{X}_{S_1}}{\partial \mathbf{H}} \right], \quad (4.4)$$

where we suppose that the price P at time S_1 is a smooth function of \mathbf{X}_{S_1} : this is true unless the payoff has a discontinuity maturing at time S_1 , in which case one can add a point to the simulation grid \mathbf{S} .

If we substitute the original SDE with its Euler discretization, in which drift and diffusion coefficients are constant on time intervals $[S_i, S_{i+1})$, we get:⁵

$$\begin{aligned} \mathbf{X}_{S_1} &= (\Sigma(\mathbf{X}_0, 0) + \mathbf{H}) \mathbf{W}_\tau + \Sigma(\mathbf{X}_0, 0) (\mathbf{W}_{S_1} - \mathbf{W}_\tau) = \Sigma(\mathbf{X}_0, 0) \mathbf{W}_{S_1} + \mathbf{H} \mathbf{W}_\tau \\ &= \Sigma(\mathbf{X}_0, 0) \mathbf{W}_{S_1} + \mathbf{H} \left(S_1^{-1} \tau \mathbf{W}_{S_1} + \sqrt{S_1^{-1} (S_1 - \tau) \tau} \tilde{\mathbf{Z}} \right) \end{aligned} \quad (4.5)$$

⁵This way we differentiate the price *as approximated by the pricing engine*; this should make little difference in practice, and is anyway arguably even more useful than the differentiation of the abstract price which would arise from the theoretical exact solution of the SDE. Note that the argument in fact needs this approximation only in the first time interval $[0, S_1)$.

for $\tilde{\mathbf{Z}} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_{N_W})$ independent from \mathbf{W}_S , as prescribed by the Brownian bridge distribution of \mathbf{W}_τ given \mathbf{W}_S (see e.g. Karatzas and Shreve, 1991). This means that we can rewrite (4.4) as

$$\frac{\partial P'}{\partial \mathbf{H}} = \mathbb{E}[\mathbf{W}_\tau \cdot \nabla P(\mathbf{X}_{S_1}, S_1)^\top] = \mathbb{E}\left[\left(\frac{\tau}{S_1} \mathbf{W}_{S_1} + \sqrt{\frac{(S_1 - \tau)\tau}{S_1}} \tilde{\mathbf{Z}}\right) \cdot \nabla P(\mathbf{X}_{S_1}, S_1)^\top\right].$$

This expression is not yet ready for ordinary differentiation with respect to τ because of the $\sqrt{\tau}$ dependence in the second addend. However, since we are only concerned with the value of the right hand side in $\mathbf{H} = \mathbf{0}$, in which point \mathbf{X}_{S_1} is a function of \mathbf{W}_{S_1} but not of $\tilde{\mathbf{Z}}$, we can use independence to deduce

$$\mathbb{E}\left[\sqrt{\frac{(S_1 - \tau)\tau}{S_1}} \tilde{\mathbf{Z}} \cdot \nabla P(\mathbf{X}_{S_1}, S_1)^\top\right] = \mathbb{E}\left[\sqrt{\frac{(S_1 - \tau)\tau}{S_1}} \tilde{\mathbf{Z}}\right] \cdot \mathbb{E}[\nabla P(\mathbf{X}_{S_1}, S_1)^\top] = \mathbf{0}.$$

To sum up, to apply (4.3), we have to estimate simply

$$\frac{\partial}{\partial \tau} \mathbb{E}\left[\frac{\tau}{S_1} \mathbf{W}_{S_1} \cdot \nabla P(\mathbf{X}_{S_1}, S_1)^\top\right] = \mathbb{E}\left[\frac{\mathbf{W}_{S_1}}{S_1} \cdot \bar{\mathbf{X}}_{S_1}\right] \quad (4.6)$$

where in the second equality we have substituted the Delta sensitivity ∇P with any estimator thereof.

The final equation (4.6) closely resembles that of likelihood-ratio based estimators like LRMAAD, while being both simpler to implement since it does not involve second order differentiation of the drift and diffusion coefficients, and more flexible since it does not prescribe which first order estimator $\bar{\mathbf{X}}_{S_1}$ to use. This also means that it shares the main drawback of likelihood ratio weights, i.e. the reciprocal dependence on the small discretization step S_1 , which may impact negatively the Monte Carlo variance. However, as is the case in vibrato, a simple antithetic trick can save the day, since the resulting estimator becomes

$$\Sigma^\top(\mathbf{X}_0, 0)^{-1} \frac{\mathbf{W}_{S_1}}{2S_1} \cdot [\bar{\mathbf{X}}_{S_1}(\mathbf{W}_{S_1}) - \bar{\mathbf{X}}_{S_1}(-\mathbf{W}_{S_1})], \quad (4.7)$$

where, roughly speaking, we expect

$$\mathbf{W}_{S_1} = O(\sqrt{S_1}) \quad \text{and} \quad \bar{\mathbf{X}}_{S_1}(\mathbf{W}_{S_1}) - \bar{\mathbf{X}}_{S_1}(-\mathbf{W}_{S_1}) = O(\mathbf{W}_{S_1}) = O(\sqrt{S_1}),$$

jointly counterbalancing the S_1^{-1} factor.

We stress that for maximum effectiveness of this variant, one should change the sign only of \mathbf{W}_{S_1} , unlike in standard antithetic Monte Carlo which inverts the full Brownian path. This does not prevent the usage of ordinary antithetic trajectories on top of the final estimator (4.7), as one would do for any other method.

5 Empirical results

In this section, we will compare empirically the above algorithms on several test cases. We will always suppose that the payoff is written on one or more underlying assets A_i with the simplest Black-Scholes dynamics

$$A_i = \exp(X_i), \quad dX_i = -\frac{\sigma_i^2}{2} dt + \sigma_i dB_i, \quad d\langle B_i, B_j \rangle = \rho_{ij} dt, \quad (5.1)$$

where B_i are correlated Brownian motions; of course, in this stylized case Euler simulation is not strictly necessary since finite-step transition probabilities are available in closed form, but we neglect this simplification which would not apply to general diffusive models. Note that (5.1) is readily expressed in the form (2.1), which formally prescribes independent Brownian motions, by a Choleski factorization of the instantaneous correlation matrix $\mathbf{R} = (\rho_{ij})_{i,j \leq N_W}$:

$$\begin{aligned} d\mathbf{X} &= \boldsymbol{\mu} dt + \boldsymbol{\Sigma} d\mathbf{W}, \\ \boldsymbol{\mu} &:= (-\sigma_i^2/2)_{i \leq N_W} \quad \text{and} \quad \boldsymbol{\Sigma} := \text{diag}(\sigma_i)_{i \leq N_W} \cdot \text{Choleski}(\mathbf{R}). \end{aligned}$$

The following names will be used to denote subsets of the full Hessian \mathbf{H} :

$$\text{Gamma} = \left(\frac{\partial^2 P}{\partial X_i(0) \partial X_j(0)} \right)_{i,j}, \quad \text{Vanna} = \left(\frac{\partial^2 P}{\partial X_i(0) \partial \sigma_j} \right)_{i,j}, \quad \text{Volga} = \left(\frac{\partial^2 P}{\partial \sigma_i \partial \sigma_j} \right)_{i,j};$$

this is slight abuse of terminology, since to adhere to common jargon one should substitute in the definitions \mathbf{X} with \mathbf{A} , but in our numerical analysis we prefer to display the sensitivity with respect to the simulated driver \mathbf{X} , because for most estimators this is more natural. One can always obtain the sensitivity to \mathbf{A} plugging the result into the elementary relations

$$\begin{aligned} \frac{\partial^2 P}{\partial A_i(0) \partial A_j(0)} &= \frac{1}{A_i(0) A_j(0)} \frac{\partial^2 P}{\partial X_i(0) \partial X_j(0)} - \delta_{ij} \frac{1}{A_i(0)^2} \frac{\partial P}{\partial X_i(0)}, \\ \frac{\partial^2 P}{\partial A_i(0) \partial \sigma_j} &= \frac{1}{A_i(0)} \frac{\partial^2 P}{\partial X_i(0) \partial \sigma_j}. \end{aligned}$$

In the sequel, all algorithms will be denoted by acronyms: we point to Tables 1 and 2 for easy reference. Note that FT can be based on any first order estimator, hence we will use the notations FT-PW, FT-VB, FT-OPP, FT-DAAD to denote the functional Gamma method of Section 4.3 where the first order estimator used to compute $\bar{\mathbf{X}}_{S_1}$ is respectively given by pathwise differentiation, vibrato Monte Carlo, Optimal Partial Proxy and DAAD.

All the test payouts will be autonomous in the sense of (2.4), so that all methods described above apply; only when the payoff is discontinuous, some of them will be unavoidably excluded. In fact in those cases one might smooth the payoff, not only to enable otherwise unusable estimators, but also to try and improve the performance of other ones. We neglect this possibility since smoothing techniques would involve a careful choice of their parameters to correctly balance the variance gain with the bias introduced, which is problematic already for first order estimators: see the analysis in Daluiso and Facchinetti (2018).

In fact, also the choice of the displacement for finite difference based algorithms (FDIFF2, FDPW, VFD, FDDAAD) may be a serious concern. We will use for $\Delta X_i(0)$ and $\Delta \sigma_i$ the largest power of 10 such that the confidence interval of the finite difference estimator contains the analytical result when available, or does not appear to diverge significantly from the confidence intervals of unbiased methods otherwise. This is a bit of cheating, and is only justified because we are just looking for benchmarks to the analytical methods which are the main subject of our analysis; however, if one would actually choose a finite difference method as his only

sensitivity calculation engine, such comparison terms would not be available, and much attention should be paid to accurate heuristics for the choice of a good finite shift.

5.1 Single-asset payoffs: stability of the estimators

Before analysing settings in which the high number of sensitivities to compute is a major driver in the choice of the algorithm, we deem it useful to explore the behaviour of the several estimators in the simpler task of computing the Gamma sensitivity with respect to a single underlying asset. This way, we can for a moment forget about the computational burden, which is comparable and affordable for all methods, and concentrate on more intrinsic and implementation-independent properties, among which we put the focus on the following:

- Variance: how large is the confidence interval for a fixed number of Monte Carlo paths? We will analyse the standard deviation of the result for 10^5 simulated paths, which by the central limit theorem should be proportional to the width of a centred confidence interval for any confidence level desired.
- Stability: does the result change smoothly while moving the initial condition S_0 ? This is a highly desirable property in practice, since one does not want the Greeks to jump wildly as the market evolves.

In all experiments, the volatility parameter is set to $\sigma = 0.2$ in yearly time units.

5.1.1 Call option

As a prototypical example of a continuous payoff we take a vanilla call option $(A_T - K)^+$ with strike $K = 100$, and vary the initial condition A_0 and the maturity time T . In particular, we take both a moderate maturity T of one year, divided for the purpose of Euler simulation into 100 equal time intervals; and a short maturity $T = 1$ day, with hourly steps in the Euler simulation. The analytical price is well known and will serve as a check.

First of all, we compare the Monte Carlo uncertainties in Table 3. We make two main remarks:

- Likelihood ratio based methods have by far the largest confidence intervals, followed by functional Gamma.
- For this simple payoff, only DAAD2 and VAD offer a material improvement over naive finite differences.

As far as smooth dependence on A_0 is concerned, we first consider the case in which T is one year. Figure 1 plots the complete results: all methods appear to work reasonably well, but those based on likelihood ratio or functional Gamma display larger errors, because of the larger confidence intervals. The smoothest analytical methods are HOPP and DAAD2; also VAD performs well, while FDPW is quite more erratic. However, one should again note that for this simple payoff, elementary second order finite differences FDIFF2, here with displacement $\Delta X_0 = 0.01$, are not that bad.

Method	$T = 1d$	$T = 1y$
FDIFF2	48.562839	3.604733
FDPW	41.621283	3.097918
HOPP	47.695944	3.719352
LRMAAD	104.264491	12.588643
AADLRM	104.264491	12.588643
VAD	23.330228	1.881182
DAAD2	19.015445	1.552514
FT-PW	53.021142	6.652639
FT-VB	65.882049	8.471453

Table 3: Gamma sensitivity of a vanilla call: standard deviation for 10^5 Monte Carlo paths, averaged over 40 values of A_0 .

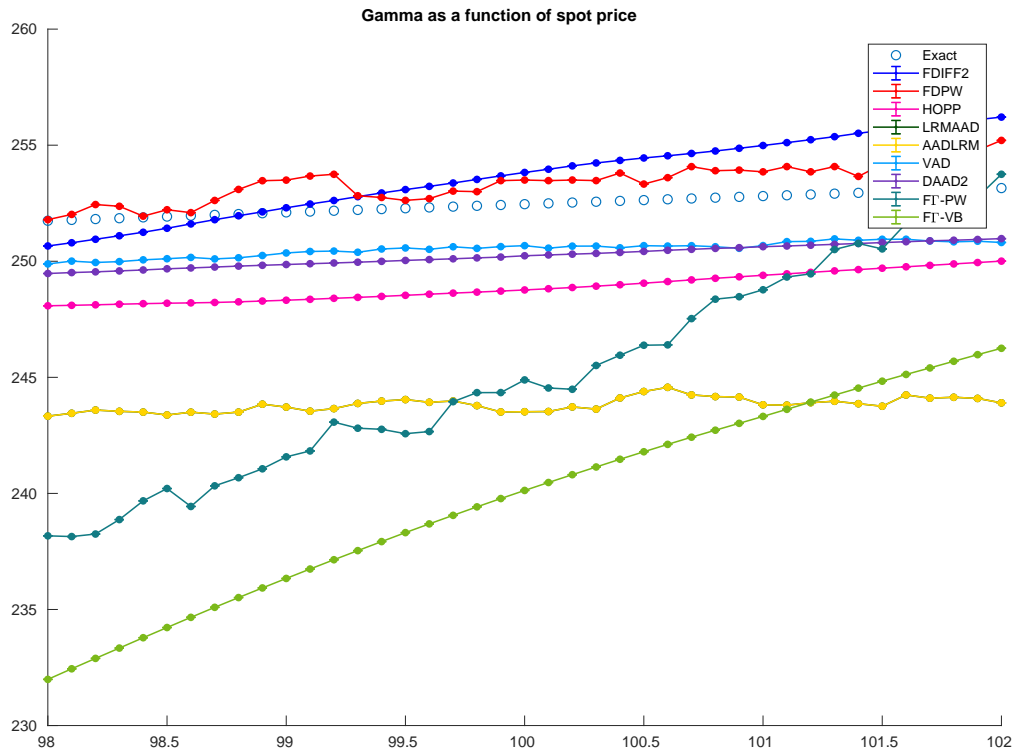


Figure 1: Gamma sensitivity of a vanilla call with maturity $T = 1y$ and strike $K = 100$ as a function of the spot price A_0 : all estimators, 10^5 paths.

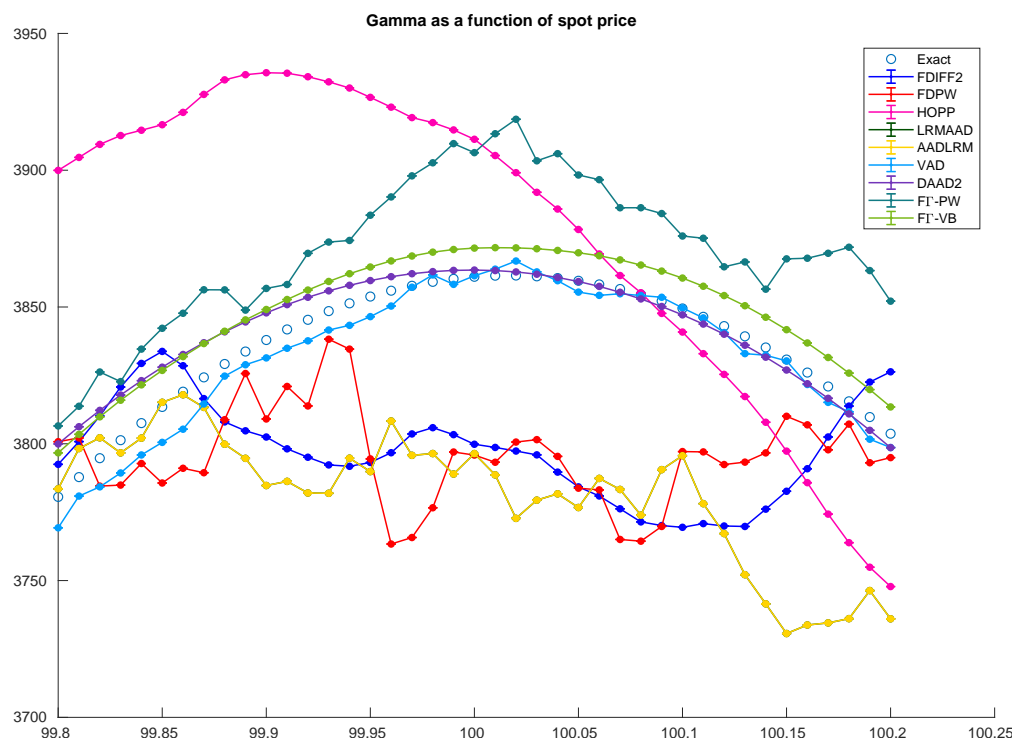


Figure 2: Gamma sensitivity of a vanilla call with maturity $T = 1d$ and strike $K = 100$ as a function of the spot price A_0 : all estimators, 10^5 paths.

We tried to stress the exercise considering an option expiring in $T = 1$ day: the results are in Figure 2. One notable difference is that FDIFF2 ($\Delta X_0 = 0.001$) and HOPP lose some smoothness, while DAAD2 still follows the shape of the exact solution remarkably well. Moreover, FT-VB becomes the most stable alternative, even though it is still significantly more noisy than most of its competitors, as we saw in Table 3. This smoothness is not observed in FT-PW, showing that the flexibility in the choice of the underlying first order estimator which functional Gamma offers over traditional likelihood ratio methods is a valuable asset.

5.1.2 Digital option

The archetypal example of a discontinuous payoff is the digital option $I_{A_T > K}$, for which we repeat the analyses of the previous subsection. Note that several methods cannot be used any more.

The detail of Monte Carlo uncertainties is in Table 4. The winners are FT-DAAD, DAAD2 and FT-VB, followed by the biased VFD; all methods are much better than FDIFF2, so we exclude it from the graphs of this subsection.

Figure 3 shows that for $T = 1$ year, DAAD based methods (DAAD2, FDDAAD, FT-DAAD) display the smoothest dependence on A_0 . The same remarks apply to $T = 1$ day, which was plotted in Figure 4, where the analytical exact Gamma has been subtracted from the estimated value for a better display cleaned from the trend.

Method	$T = 1d$	$T = 1y$
FDIFF2	865.468515	6.303787
HOPP	198.053629	1.672700
VFD	88.642863	0.699658
FDDAAD	134.068248	1.095060
DAAD2	73.659168	0.587246
FT-VB	75.564853	0.652947
FT-OPP	123.806197	1.030655
FT-DAAD	61.683419	0.547219

Table 4: Gamma sensitivity of a digital call: standard deviation for 10^5 Monte Carlo paths, averaged over 40 values of A_0 .




Figure 3: Gamma sensitivity of a digital call with maturity $T = 1y$ and strike $K = 100$ as a function of the spot price A_0 : low-variance estimators, 10^5 paths.

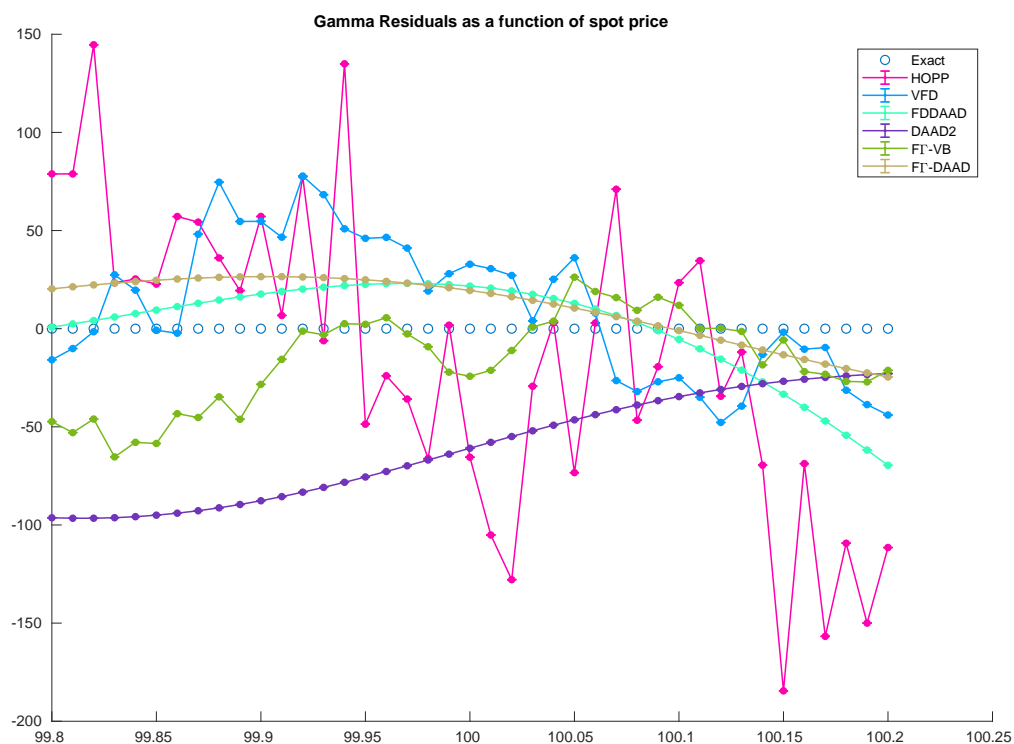


Figure 4: Gamma sensitivity of a digital call with maturity $T = 1d$ and strike $K = 100$ as a function of the spot price A_0 : residual of low-variance estimators with respect to the analytical value, 10^5 paths.

5.2 Multi-asset payoffs: efficiency of the algorithms

We now move to multi-asset payoffs, where the full Hessian \mathbf{H} has a high number of entries. This imposes a rethinking of how the various estimators are compared, because a more noisy one may still be preferable if it is significantly faster. We believe that the most objective metric is therefore the estimated run time which would be needed to achieve a unit Monte Carlo variance. This is easily computed as

$$(\text{run time}) \times (\text{estimated standard deviation})^2, \quad (5.2)$$

since it is well known that to reduce the uncertainty by a factor of two one should increase the number of simulated paths by a factor of four.

In all examples, all assets have spot value $A_i(0) = 100$ and volatility parameter $\sigma_i = 0.2$, while the instantaneous correlation between couples of Brownian drivers is set to a flat value $\rho_{ij} = 0.5$. Each method was tested on a Monte Carlo simulation with 10^4 paths.

We will plot the chosen metric (5.2) separately for every second order sensitivity in the Hessian matrix, because while on the one hand the run time is a common number, on the other hand the standard deviation is different entry by entry. In all the figures, the graphs on the left concern the sensitivities in which both differentiation variables refer to the same asset, while all other entries are plotted in the graphs on the right.

5.2.1 Basket call

The multidimensional generalization of the first example of Section 5.1 is a call on the average of several assets:

$$\left(\frac{1}{N_X} \sum_{i=1}^{N_X} A_i(T) - K \right)^+,$$

which we will price for $N_X = 8$, $T = 1$ year and $K = 100$.

The results are in Figures 5 and 6; note the absence of likelihood ratio methods, since their values are between one and two orders of magnitude larger than those of the other algorithms, and the graph would be unreadable. In particular, finite differences should be preferred to likelihood ratio regardless of their higher cost per path; while all other methods outperform FDIFF2, as the plot shows. DAAD2 is consistently much more effective than its competitors, while the second-best depends on the sensitivity of interest: it is VAD for Gamma, but becomes FDPW for Volga, VAD for Delta of Vega, HOPP for Vega of Delta (although Delta of Vega and Vega of Delta estimate theoretically the same number by the Schwarz theorem, see the last paragraph of this section).

If one wants to compute only Gamma, then functional methods are also available; moreover, LRMAAD might in principle become competitive thanks to its efficient implementation of Section 3.2. For a fair comparison we updated also the run times of the other methods, which are smaller if sensitivities to σ_i are not required. Our empirical findings in this setting are in Figure 7, and are not encouraging: LRMAAD is still far too noisy even considering that it is much faster than

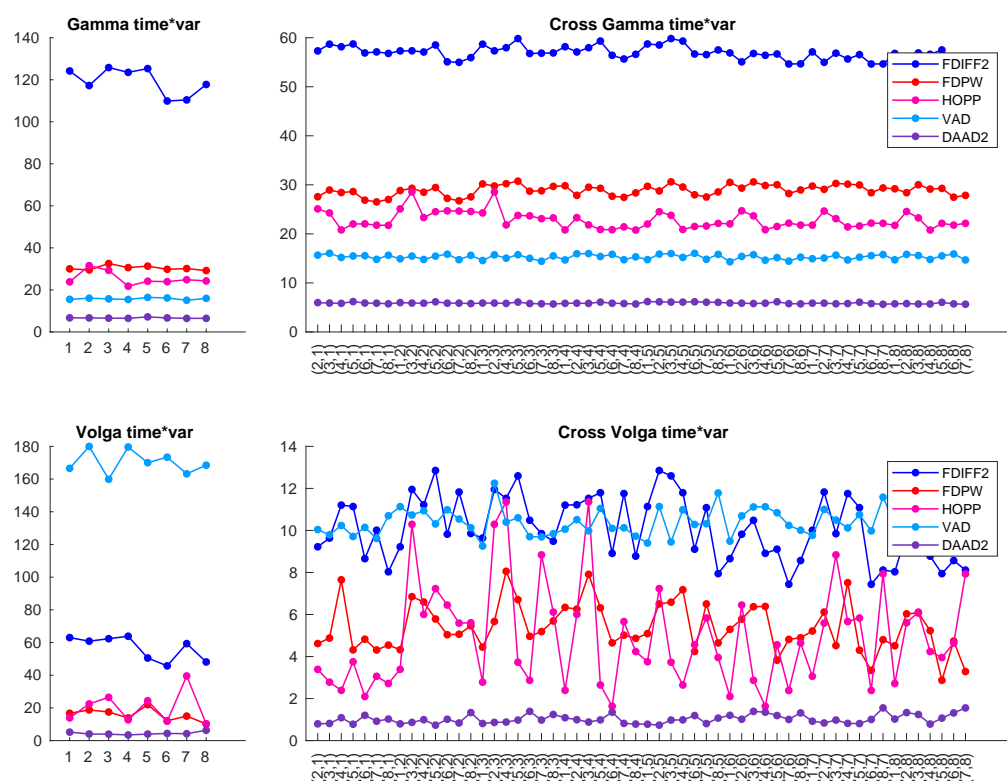


Figure 5: Gamma and Volga sensitivities of a call with maturity $T = 1y$ and strike $K = 100$ on a basket of 8 underlyings: efficiency of all methods except likelihood ratio (out of scale).

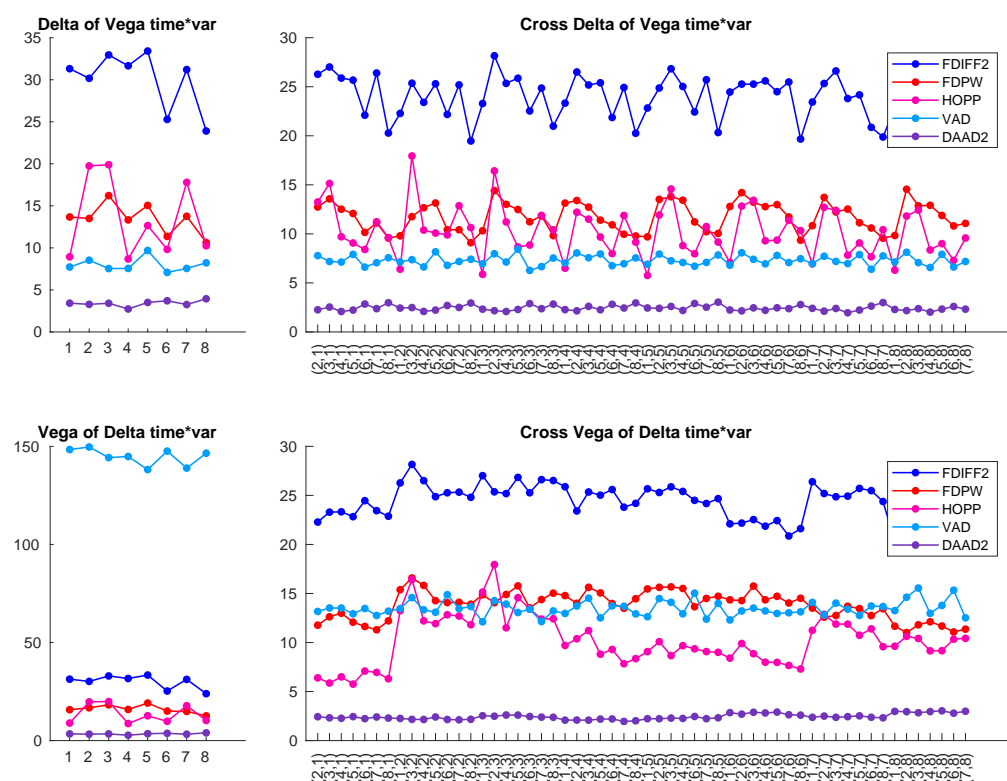


Figure 6: Vanna sensitivities of a call with maturity $T = 1y$ and strike $K = 100$ on a basket of 8 underlyings: efficiency of all methods except likelihood ratio (out of scale).




Figure 7: Gamma sensitivities of a call with maturity $T = 1y$ and strike $K = 100$ on a basket of 8 underlyings: efficiency of all methods except AADLRM (out of scale).




Figure 8: Gamma sensitivities of a call with maturity $T = 1y$ and strike $K = 100$ on a basket of 8 underlyings: efficiency of all methods except likelihood ratio and functional Gamma (out of scale).

linear-cost methods (not to speak of FDIFF2), and FT, while significantly better, does not improve the variance sufficiently. We remark that now that the number of required sensitivities is “only” 64, beating the dumb FDIFF2 benchmark is harder than one might expect, and off the diagonal only DAAD2 is a neat improvement (Figure 8, right-hand graph).

Finally, before accepting this disappointing remark, we note that most of the analytical estimators we test produce asymmetric matrices, although we know that Hessians are symmetric: so we have two distinct estimators for each off-diagonal entry. Therefore, we can look for a linear combination to reduce the Monte Carlo variance: equivalently, we can correct the estimator Γ_{ij} using the mean-zero $(\Gamma_{ji} - \Gamma_{ij})$ as a control variate. The coefficients will be chosen by a small-sample preliminary Monte Carlo run as is customary for control variates (Glasserman, 2004). Some methods benefit from this trick more than others, as one can see from Figure 9: in particular, functional Gamma comes closer to finite differences (but is still more than twice slower in getting unit variance), and VAD becomes as good as DAAD2 on Cross Gammas.




Figure 9: Gamma sensitivities of a call with maturity $T = 1y$ and strike $K = 100$ on a basket of 8 underlyings: efficiency of FDIFF2 and of symmetrized analytical estimators except AADLRM (out of scale).

5.2.2 Basket digital

A good representative example of a discontinuous multi-asset product is a digital option on a basket:

$$I_{(K,+\infty)} \left(\frac{1}{N_X} \sum_{i=1}^{N_X} A_i(T) \right).$$

As in the previous subsection, we take $N_X = 8$, $T = 1$ year and $K = 100$.

The results for the full Hessian are in Figures 10 and 11: they show that FDIFF2 is inadequate for discontinuous multi-asset payoffs, while FDDAAD and HOPP are not optimal either. DAAD2 and VFD are comparable.

The biggest surprise, however, comes from the estimators restricted to the Gamma matrix in Figure 12. Indeed, we see that the best linear-cost methods DAAD2 and VFD are clearly overcome by constant-cost functional Gamma methods, with FT-DAAD being 4-5 times faster than the fastest linear-complexity algorithm (always in uncertainty-adjusted terms). This is at odds with the results obtained in the previous subsection on the “easy” call payoff, probably because on continuous payouts traditional methods already work very well.

6 Conclusions

In this paper, we have compared a wide range of old and new methods for the computation of second order sensitivities of Monte Carlo prices of financial products, with special emphasis on algorithmic efficiency when a large number of Greeks must be computed.

From the literature review viewpoint, the theoretical analysis has examined, with extensions (LRMAAD, VAD) and variations (VFD), the main algorithms to compute the full Hessian matrix of an expected value when the underlying is driven by a multidimensional Brownian diffusion. Several limitations of applicability are found:

1. Many methods often have too high Monte Carlo uncertainties: in particular,




Figure 10: Gamma and Volga sensitivities of a digital call with maturity $T = 1y$ and strike $K = 100$ on a basket of 8 underlyings: efficiency of all methods.

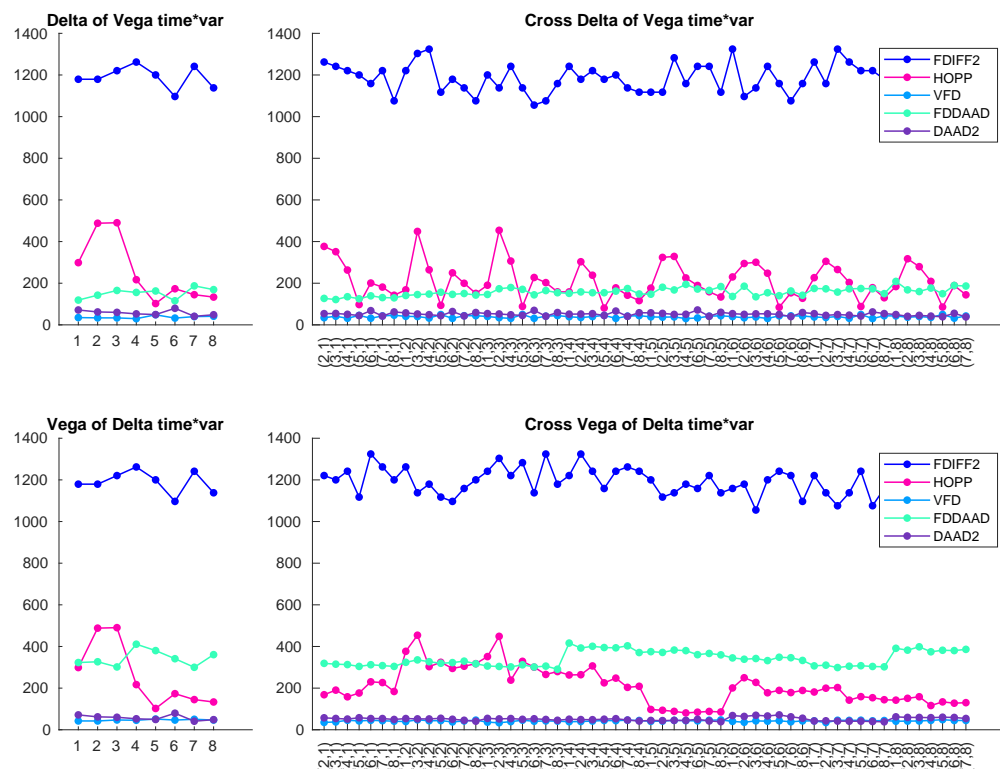


Figure 11: Vanna sensitivities of a digital call with maturity $T = 1y$ and strike $K = 100$ on a basket of 8 underlyings: efficiency of all methods.




Figure 12: Gamma sensitivities of a digital call with maturity $T = 1y$ and strike $K = 100$ on a basket of 8 underlyings: efficiency of best-performing methods.

those based on likelihood ratio (LRMAAD, AADLRM), and to a lesser extent those based on finite differences of the pathwise estimator (FDPW).

2. Most methods only apply when there is at least one random driver per underlying ($N_X \leq N_W$) and the payoff does not depend directly on the initial state \mathbf{X}_0 , nor on the model parameters $\boldsymbol{\theta}$ (the exceptions being HOPP and the dumb FDPW).
3. Most methods cannot handle discontinuous payoffs (the exceptions being HOPP and VFD).
4. For all methods except LRMAAD restricted to sensitivities to \mathbf{X}_0 only, the multiplicative overhead on the run time for only-price grows linearly in the number of rows of the Hessian.

The original estimators of this paper should be considered in view of the above limitations of their many competitors (see Table 5 for a visual summary):

- DAAD2 is the only method besides HOPP which solves simultaneously issues 1, 2 and 3, and in our tests was always significantly more effective than the latter from the statistical uncertainty point of view, while showing remarkable smoothness properties with respect to small changes in the input parameters.
- FT is the first estimator solving simultaneously issues 1 and 4 in at least one relevant case; it is also the first constant-complexity algorithm applicable to discontinuous payoffs. Its best-of-class efficacy in the basket digital test case, coupled with the easy implementation, qualify it as a must-try when tackling a new tough high-dimensional problem. Since it is in fact a meta-algorithm based on any first order estimator, it offers significant flexibility for either seamless introduction in any financial library already capable of computing first order sensitivities, or for fine tuning to specific payoffs.

Besides systematic testing on a large range of payoff types and model dynamics, a couple of issues remain open for further research. In the field of linear-cost methods, the main challenge is related to payouts with many discontinuities, for which the computational complexity of the adjustments becomes relevant, particularly when it scales quadratically in the number of digital features as for the current implementation of DAAD2. Moreover, this paper opens the new field of constant-cost methods, where a general-purpose algorithm is still to be found, since FT in one example showed poor performance, and is anyway limited to derivatives with respect to \mathbf{X}_0 . Finally, one might consider the possibly varied effectiveness of variance reduction techniques to different estimators.

7 Acknowledgements

The author is grateful to Giorgio Facchinetti and Massimo Morini for their comments on the drafts of this paper, and to Giulio Sartorelli, Giuseppe Di Poto and Andrea Pallavicini for helpful discussions on the topic.

Table 5: Properties and domain of applicability of the methods.

	FDPW	LRMAAD	AADLRM	VAD	HOPP	DAAD2	FF
Low variance	✗	✗	✗	✓	✓	✓	✗/✓ ¹
Non-autonomous	✓	✗	✗	✗	✓	✓	✓
Discontinuous	✗	✗	✗	✓ ²	✓	✓	✓ ³
Constant-cost Γ	✗	✓	✗	✗	✗	✗	✓
Full Hessian	✓	✓	✓	✓	✓	✓	✗

¹ Depending on the payoff.

² In fact one must substitute AD with a finite difference (VFD).

³ If handled by the underlying algorithm.

References

- Basel Committee on Banking Supervision (2015). Review of the credit valuation adjustment risk framework.
- Capriotti, L. (2011). Fast Greeks by algorithmic differentiation. *Journal of Computational Finance*, 14(3):3–35.
- Capriotti, L. (2015). Likelihood ratio method and algorithmic differentiation: Fast second order Greeks. *Algorithmic Finance*, 4:81–87.
- Capriotti, L. and Giles, M. (2012). Adjoint Greeks made easy. *Risk*, (September):92–99.
- Chan, J. H. and Joshi, M. (2015). Optimal limit methods for computing sensitivities of discontinuous integrals including triggerable derivative securities. *IEEE Transactions*, 47:978–997.
- Daluiso, R. and Facchinetti, G. (2018). Algorithmic differentiation for discontinuous payoffs. *International Journal of Theoretical and Applied Finance*, 21.
- Giles, M. (2007). Monte Carlo evaluation of sensitivities in computational finance. Technical Report NA07/12, Oxford University Computing Lab.
- Giles, M. (2009). Vibrato Monte Carlo sensitivities. In L’Ecuyer, P. and Owen, A. B., editors, *Monte Carlo and Quasi-Monte Carlo Methods 2008*, pages 369–382, Berlin. Springer.
- Glasserman, P. (2004). *Monte Carlo Methods in Financial Engineering*. Springer, Berlin.
- Griewank, A. and Walther, A. (2008). *Evaluating Derivatives: Principles and Techniques of Algorithmic Differentiation*. Society for Industrial and Applied Mathematics, Philadelphia, second edition.
- Joshi, M. and Yang, C. (2011). Algorithmic Hessians and the fast computation of cross-gamma risk. *IEEE Transactions*, 43:878–892.

- Joshi, M. and Zhu, D. (2016). Optimal partial proxy method for computing gammas of financial products with discontinuous and angular payoffs. *Applied Mathematical Finance*, 23(1):22–56.
- Karatzas, I. and Shreve, S. E. (1991). *Brownian Motion and Stochastic Calculus*. Springer, New York.
- Pagès, G., Pironneau, O., and Sall, G. (2016). Vibrato and automatic differentiation for high order derivatives and sensitivities of financial options. arXiv:1606.06143v1.
- Reghai, A., Kettani, O., and Messaoud, M. (2015). CVA with Greeks and AAD. *Risk*, (December).

A Second order distributional pathwise differentiation

In this Appendix, we will derive an estimator for the Hessian of the price

$$P = \mathbb{E} [g(\mathbf{X}_T, \boldsymbol{\theta})]$$

for a piecewise smooth but possibly discontinuous payoff g , by a generalization of the Dirac-delta based method DAAD of Daluiso and Facchinetti (2018), from which we take the following conventions: given a vector $\mathbf{v} \in \mathbb{R}^d$ and $k \in \{1, \dots, d\}$, we will denote by $\mathbf{v}_{-k} \in \mathbb{R}^{d-1}$ the vector obtained from \mathbf{v} removing its k -th component, and with $\mathbf{v}(v_k = x)$ the vector obtained from \mathbf{v} substituting the k -th component with the value x ; we will denote by $\mathbf{I}_{\mathbf{v} > \mathbf{0}}$ the vector $(I_{v_1 > 0}, \dots, I_{v_d > 0})$; finally, for a function ψ defined on $\{0, 1\}^h$ and for $i = 1, \dots, h$, we define

$$\Delta^{(i)}\psi(\mathbf{a}) = \psi(\mathbf{a}(a_i = 1)) - \psi(\mathbf{a}(a_i = 0)).$$

In view of (4.1), we just need to differentiate twice the conditional expectation ε of Section 4.2, which has the following form:

$$\varepsilon = \mathbb{E} [\phi(I_{f_1(\boldsymbol{\theta}, \mathbf{Z}) > 0}, \dots, I_{f_h(\boldsymbol{\theta}, \mathbf{Z}) > 0}, \boldsymbol{\theta}, \mathbf{Z})] = \mathbb{E} [\phi(\mathbf{I}_{\mathbf{f}(\boldsymbol{\theta}, \mathbf{Z}) > \mathbf{0}}, \boldsymbol{\theta}, \mathbf{Z})],$$

where \mathbf{Z} is a standard Gaussian random vector, and ϕ and the f_i are suitable functions, smooth in $\boldsymbol{\theta}$ and \mathbf{Z} . We make the following weak assumption:

Hypothesis A.1. f_i is twice differentiable with respect to $\boldsymbol{\theta}$ and $z_{k(i)}$. Moreover, for almost every $\mathbf{Z}_{-k(i)}$, the function $z_{k(i)} \mapsto f(\boldsymbol{\theta}, \mathbf{Z}(Z_{k(i)} = z_{k(i)}))$ is zero only for $z_{k(i)}$ in a finite set $A_{k(i)}^i(\boldsymbol{\theta}, \mathbf{Z}_{-k(i)})$, and its derivative in such points is non-null.

We also suppose for notational convenience and with no loss of generality that

$$\Delta^{(i)}\phi(\mathbf{I}_{\mathbf{f}(\boldsymbol{\theta}, \mathbf{Z}) > \mathbf{0}}, \boldsymbol{\theta}, \mathbf{Z}) = 0 \quad \forall i < c,$$

so that the indicator functions before c represent the discontinuities in the gradient of the payoff, while the indicator functions from c onwards represent the discontinuities of the payoff itself, if any.

The cited paper tells us that the first order sensitivity is

$$\frac{\partial \varepsilon}{\partial \boldsymbol{\theta}} = \mathbf{d}_0 + \sum_{i=c}^h \mathbf{d}_i,$$

where

$$\begin{aligned} \mathbf{d}_0 &= \mathbb{E} \left[\frac{\partial \phi}{\partial \boldsymbol{\theta}} \left(\mathbf{I}_{\mathbf{f}(\boldsymbol{\theta}, \mathbf{Z}) > \mathbf{0}}, \boldsymbol{\theta}, \mathbf{Z} \right) \right], \\ \mathbf{d}_i &= \mathbb{E} \left[\sum_{\zeta \in A_{k(i)}^i(\boldsymbol{\theta}, \mathbf{Z}_{-k(i)})} \left\{ \frac{e^{-\zeta^2/2}}{\sqrt{2\pi}} \left[\left| \frac{\partial f_i}{\partial z_{k(i)}} \right|^{-1} \frac{\partial f_i}{\partial \boldsymbol{\theta}} \right] (\boldsymbol{\theta}, \mathbf{Z}(Z_{k(i)} = \zeta)) \right. \right. \\ &\quad \left. \left. \times \Delta^{(i)} \phi \left(\mathbf{I}_{\mathbf{f}(\boldsymbol{\theta}, \mathbf{Z}(Z_{k(i)} = \zeta)) > \mathbf{0}}, \boldsymbol{\theta}, \mathbf{Z}(Z_{k(i)} = \zeta) \right) \right\} \right]. \end{aligned}$$

Now we have to differentiate \mathbf{d}_0 and \mathbf{d}_i . The first task is a direct application of the first order result to $\nabla \phi$:

$$\frac{\partial \mathbf{d}_0}{\partial \boldsymbol{\theta}} = \mathbf{H}_{00} + \sum_{i=1}^h \mathbf{H}_{0i},$$

where

$$\begin{aligned} \mathbf{H}_{00} &= \mathbb{E} \left[\frac{\partial^2 \phi}{\partial \boldsymbol{\theta}^2} \left(\mathbf{I}_{\mathbf{f}(\boldsymbol{\theta}, \mathbf{Z}(Z_{k(i)} = \zeta)) > \mathbf{0}}, \boldsymbol{\theta}, \mathbf{Z} \right) \right], \\ \mathbf{H}_{0i} &= \mathbb{E} \left[\sum_{\zeta \in A_{k(i)}^i(\boldsymbol{\theta}, \mathbf{Z}_{-k(i)})} \left\{ \Delta^{(i)} \left(\frac{\partial \phi}{\partial \boldsymbol{\theta}} \right)^\top \left(\mathbf{I}_{\mathbf{f}(\boldsymbol{\theta}, \mathbf{Z}(Z_{k(i)} = \zeta)) > \mathbf{0}}, \boldsymbol{\theta}, \mathbf{Z}(Z_{k(i)} = \zeta) \right) \right. \right. \\ &\quad \left. \left. \times \frac{e^{-\zeta^2/2}}{\sqrt{2\pi}} \left[\left| \frac{\partial f_i}{\partial z_{k(i)}} \right|^{-1} \frac{\partial f_i}{\partial \boldsymbol{\theta}} \right] (\boldsymbol{\theta}, \mathbf{Z}(Z_{k(i)} = \zeta)) \right\} \right]. \end{aligned}$$

If the payoff is continuous ($c = h + 1$), we are done; otherwise, we need the derivative of \mathbf{d}_i , which is slightly trickier. In order to compute it, we note that by the implicit function theorem, in a neighbourhood of $\boldsymbol{\theta}$, the zeros ζ in the definition of each \mathbf{d}_i can be expressed as functions $\zeta_i^{(l)}$ of $\mathbf{Z}_{-k(i)}$ and $\boldsymbol{\theta}$ with gradient⁶

$$- \left[\left(\frac{\partial f_i}{\partial z_{k(i)}} \right)^{-1} \left(\frac{\partial f_i}{\partial (\boldsymbol{\theta}, \mathbf{z}_{-k(i)})} \right) \right] (\boldsymbol{\theta}, \mathbf{Z}(Z_{k(i)} = \zeta)).$$

Now each summand in the definition of \mathbf{d}_i can be differentiated with the DAAD method, provided that the composition of f_j with $\zeta_i^{(l)}$ satisfies Hypothesis A.1 for some coordinate $h(i, j) \neq k(i)$. We define

$$A^{i,j}(\boldsymbol{\theta}, \mathbf{Z}) := \{(\zeta, \eta) : (f_i, f_j)(\boldsymbol{\theta}, \mathbf{Z}(Z_{k(i)} = \zeta, Z_{h(i,j)} = \eta)) = (0, 0)\}.$$

⁶In fact the neighbourhood may depend on \mathbf{Z} ; we neglect this kind of technicalities, which should be addressable along the lines of the Appendix of Daluiso and Facchinetti (2018).

We are ready to state the result:

$$\frac{\partial \mathbf{d}_i}{\partial \boldsymbol{\theta}} = \mathbf{H}_{i0} + \sum_{j=c}^h \mathbf{H}_{ij},$$

where

$$\begin{aligned} \mathbf{H}_{i0} &= \mathbb{E} \left\{ \sum_{\zeta \in A_{k(i)}^i(\boldsymbol{\theta}, \mathbf{Z}_{-k(i)})} \frac{\partial}{\partial(\boldsymbol{\theta}, z_{k(i)})} \Big|_{z_{k(i)}=\zeta} \left[\frac{e^{-\zeta^2/2}}{\sqrt{2\pi}} \left| \frac{\partial f_i}{\partial z_{k(i)}} \right|^{-1} \left(\frac{\partial f_i}{\partial \boldsymbol{\theta}} \right)^\top \Delta^{(i)} \phi \right] \right. \\ &\quad \times \left. \left(- \left[\left(\frac{\partial f_i}{\partial z_{k(i)}} \right)^{-1} \left(\frac{\partial f_i}{\partial \boldsymbol{\theta}} \right) \right] (\boldsymbol{\theta}, \mathbf{Z}(Z_{k(i)} = \zeta)) \right) \right\}, \\ \mathbf{H}_{ij} &= \mathbb{E} \left\{ \sum_{(\zeta, \eta) \in A^{i,j}(\boldsymbol{\theta}, \mathbf{Z})} \left[\frac{e^{-\frac{\zeta^2 + \eta^2}{2}}}{2\pi} \left| \frac{\partial f_i}{\partial z_{k(i)}} \right|^{-1} \left(\frac{\partial f_i}{\partial \boldsymbol{\theta}} \right)^\top \left| \frac{\partial f_j}{\partial z_{h(i,j)}} - \frac{\partial f_j}{\partial z_{k(i)}} \left(\frac{\partial f_i}{\partial z_{k(i)}} \right)^{-1} \frac{\partial f_i}{\partial z_{h(i,j)}} \right|^{-1} \right. \right. \\ &\quad \left. \left. \left(\frac{\partial f_j}{\partial \boldsymbol{\theta}} - \frac{\partial f_j}{\partial z_{k(i)}} \left(\frac{\partial f_i}{\partial z_{k(i)}} \right)^{-1} \frac{\partial f_i}{\partial \boldsymbol{\theta}} \right) \Delta^{(j)} \Delta^{(i)} \phi(\mathbf{I}_{\mathbf{f} > \mathbf{0}}, \boldsymbol{\theta}, \mathbf{Z}) \right] \right\}_{Z_{k(i)}=\zeta, Z_{h(i,j)}=\eta}. \end{aligned}$$

Note that the expression in the second absolute value should be different from zero, or to say it differently, $\partial(f_i, f_j)/\partial(z_{k(i)}, z_{h(i,j)})$ should be full rank.

All the gradients appearing in these formulas are readily computed by algorithmic differentiation; the only practical concern is the fact that there are (up to) $(h - c + 1)^2$ terms \mathbf{H}_{ij} . In fact in most applications, f_i and f_j depend on non-overlapping sets of fixing times, so that

$$\begin{aligned} \left| \frac{\partial f_j}{\partial z_{h(i,j)}} - \frac{\partial f_j}{\partial z_{k(i)}} \left(\frac{\partial f_i}{\partial z_{k(i)}} \right)^{-1} \frac{\partial f_i}{\partial z_{h(i,j)}} \right|^{-1} \left(\frac{\partial f_j}{\partial \boldsymbol{\theta}} - \frac{\partial f_j}{\partial z_{k(i)}} \left(\frac{\partial f_i}{\partial z_{k(i)}} \right)^{-1} \frac{\partial f_i}{\partial \boldsymbol{\theta}} \right) \\ = \left| \frac{\partial f_j}{\partial z_{k(j)}} \right|^{-1} \left(\frac{\partial f_j}{\partial \boldsymbol{\theta}} \right), \end{aligned}$$

which does not depend on i any more. As a consequence, the only portion of the algorithm which truly grows quadratically with the number of payoff discontinuities is $\Delta^{(j)} \Delta^{(i)} \phi$. This should not be dramatic in the economy of a Monte Carlo pricing exercise inclusive of a costly Euler simulation, unless there is a really large number of digital features as in frequently monitored barrier options, which anyway are often already handled by regularizing the payoff accounting for touch probabilities (see e.g. Glasserman, 2004).

For ease of reference, we collect as a conclusion the full set of addends contributing to the estimator of the Hessian:

$$\frac{\partial^2 \varepsilon}{\partial \boldsymbol{\theta}^2} = \mathbf{H}_{00} + \sum_{i=1}^h \mathbf{H}_{0i} + \sum_{i=c}^h \mathbf{H}_{i0} + \sum_{i,j=c}^h \mathbf{H}_{ij}.$$